

Peter (Jonghyeok) Yoon

jonghyeokyon19@gmail.com | [linkedin.com/in/jhyoon19](https://www.linkedin.com/in/jhyoon19) | San Jose, CA | (619) 772-0251 | github.com/peterjhyoon

WORK EXPERIENCE

SAP - Business Data Cloud

Software Engineer II | Mar 2026 - Present

Palo Alto, CA

- Co-leading the system design and architecture of an in-house data lineage solution using Kafka and Temporal, enabling event-driven Spark job triggering to improve pipeline observability and data freshness of data products
- Leading the design and development of an engine-agnostic ETL framework supporting multiple data processing engines, improving pipeline reuse and reducing compute costs

Software Engineer | Jul 2025 - Mar 2026

Palo Alto, CA

- Architected a generic spark-based ETL execution framework, establishing a unified runtime to power intelligent applications across SAP lines of business while improving reusability and extensibility of data engineering workflows
- Achieved up to a 98% reduction in staging data ingestion time, from over 8 hours to 11 minutes, by migrating JSONL-based data ingestion pipeline from Structured Streaming to Batch processing and optimizing partitioning logic
- Implemented core Spark components for the first cross-tenant derived data product prototype, collaborating with stakeholders to allow aggregating data across multiple source-system tenant containers to power SAP intelligent applications

Software Engineer Intern | Jan 2025 - Jun 2025

Newport Beach, CA

- Designed and implemented a generic data ingestion pipeline, unifying ingestion across 2 different data acquisition services in SAP Business Data Cloud with Data Protection Policy and GDPR compliance, reducing developer maintenance by 50%
- Developed an LLM-driven data generation tool PoC using Python FastAPI, React and Tailwind that dynamically produces millions of synthetic data entries in minutes for performance testing; presented to the wider engineering org for adoption
- Streamlined large-scale data transformation workflows in SAP Business Data Cloud by designing and developing reusable components of a PySpark-based plugin framework, leveraging state management and integrated into production workflows
- Improved developer productivity by helping automate linting and add integration tests for 95% code coverage on spark jobs

UC San Diego IT Services

DevOps Engineer | Jan 2024 - Dec 2024

San Diego, CA

- Reduced student major data provisioning times by 35% through ETL data processing with Python and Apache Airflow
- Improved Grouper logging resource deployment times by 85%, from 35 minutes to 5 minutes, by migrating manually created AWS resources (EC2, S3, EFS, IAM, Security Group) for the Splunk logging architecture to Terraform deployment
- Streamlined Grouper AWS infrastructure configuration by constructing a Terraform CI/CD pipeline with Atlantis

RapidSOS

Software Engineer Intern | Jun 2024 - Sep 2024

New York, NY

- Reduced deployment costs by 64% by migrating CI/CD artifact triggers of the Emergency Data Exchange (EDX) code release pipelines from polling to event-based mechanisms and alleviating year-old technical debt to save thousands annually
- Developed a serverless event-driven system to track prolonged Emergency Communication Center (ECC) websocket disconnections using AWS Lambda, DynamoDB and S3 with the Python SDK and Docker

Kewtea Korea

Software Engineer Intern | May 2023 - Aug 2023

Seoul, South Korea

- Decreased monthly overhead costs by 32% and halved processing times by building automated workflows and a Python FastAPI automation system with MySQL on GCP, featuring an ML-powered ChatBot for task updates

TECHNICAL SKILLS

Languages: Python, Java, SQL (MySQL, PostgreSQL), C/C++, JavaScript, CUDA, ARM Assembly, MATLAB

Backend/Data/AI: Apache Spark, Airflow, Kafka, AWS, Terraform, LangGraph, LangChain, FastAPI, Docker, Kubernetes, Git, Bash

Frontend: React.js, React Native, Tailwind CSS, HTML/CSS, Bootstrap, Jest

EDUCATION

University of California, San Diego

B.S. Mathematics-Computer Science

San Diego, CA